

Bitcoin Mining

CS 5158/6058 Data Security and Privacy

Spring 2018

Instructor: Boyang Wang

Proof of Work

- Alice wants to prove to Bob that she did certain amount of works (e.g., no. of operations or time)
 - Bob gives Alice a [target](#)
 - Alice finds a [solution](#) for this target
 - Bob verifies this solution based on this target
- Alice needs certain time to find a solution (time consuming, but still possible); With a solution and a target, it is very easy for Bob to verify.

Proof of Work

- Alice has a message m , wants to give it to Bob
- Bob tells Alice to find a nonce n s.t.
 - $H(m||n) \leq \text{target}$
 - Otherwise he will not accept message m
- Enumerating all possible nonces is the only way to for Alice find a solution
 - Since hash function is hard to invert

Example

- Bob tells Alice to find a nonce n s.t.
 - $H(m||n) \leq x/2$;
- Example: Given $m = uc$, Alice chooses $n = a$
 - $H(uc||a) > x/2$, a is not a solution;
- Alice tries $n = b$
 - $H(uc||b) > x/2$, b is not a solution;
- Alice tries $n = c$
 - $H(uc||c) \leq x/2$, c is a solution, return

H(*)	000	001	010	011	100	101	110	111
------	-----	-----	-----	-----	-----	-----	-----	-----

- Difficulty: the no. of starting 0 bits in a target
- The no. of hash operations at Alice increases **exponentially** with difficulty

Target	Difficulty	Hashes at Alice	Hashes at Bob
$x=111$	0	1	1
$x/2=011$	1	2	1
$x/4=001$	2	4	1
$x/8=000$	3	8	1

Proof of Work

- Given difficulty d , target $t = x/2^d = 2^{L-d}$
- For message m and nonce n ,
 - $\Pr[H(m||n) \leq t] = 1/2^d$
- Average 2^d hashes to find a solution
- Example: difficulty $d = 8$, $x = 2^{20}$, $L = 20$
 - target $t = x/2^d = 2^{20}/2^8 = 2^{12}$
 - 00000 000 11 11111 11111
 - Average $2^d = 256$ hashes to find solution

Proof of Work

- Given difficulty d , target $t = 2^{L-d}$
- On average, 2^d hashes to find a solution
- SHA256 on a laptop: 200,000 hashes/second

Difficulty	No. of Hashes	Time
20	1,048,576	5.2 secs
22	4,194,304	21 secs
24	16,777,216	84 secs
26	67,108,864	?? secs
28	268,435,456	?? secs

Proof of Work

- Given difficulty d , target $t = 2^{L-d}$
- On average, 2^d hashes to find a solution
- SHA256 on a laptop: 200,000 hashes/second

Difficulty	No. of Hashes	Time
20	1,048,576	5.2 secs
22	4,194,304	21 secs
24	16,777,216	84 secs
26	67,108,864	335 secs
28	268,435,456	1,342 secs

Proof of Work

- Given difficulty d , target $t = 2^{L-d}$
- On average, 2^d hashes to find a solution
- Practice: increase to 400,000 hashes/second

Difficulty	No. of Hashes	Time (200,000)	Time (400,000)
20	1,048,576	5.2 secs	?? secs
22	4,194,304	21 secs	?? secs

Proof of Work

- Given difficulty d , target $t = 2^{L-d}$
- On average, 2^d hashes to find a solution
- Practice: increase to 400,000 hashes/second

Difficulty	No. of Hashes	Time (200,000)	Time (400,000)
20	1,048,576	5.2 secs	2.6 secs
22	4,194,304	21 secs	10.5 secs

Applications of PoW

- Defer Denial of Service Attacks:
 - Limit the no. of service requests that an attacker can submit
- Bitcoin Mining:
 - Decide which miner will add the next block to the current blockchain

DDoS Attack

- Basic setting: a server can only take a certain number of requests in a short time
 - Due to limited memory, bandwidth, etc.
- Denial-of-Service (DDoS) Attack
 - Attacker sends a **large** no. of legitimate service requests to a server in a **short** time
 - A server **cannot** handle those requests in a short time, server goes down...

Example of DDoS Attack

- Example: If UC server receives more than 30 requests within 1 min from a user/attacker, then UC server will go down
- Legitimate user Alice submits 5 requests every min
- Attacker submits 40 requests within 1 min
- UC server goes down due to this DDoS attack

Defer DDoS with PoW

- Use PoW to defer Denial-of-Service (DDoS)
 - Server uses a counter c to record the number of request that a user/attacker has submitted
 - Server uses counter c as difficulty d
 - Server sets target as $t = 2^{L-c}$
 - For each new request r , a user/attacker is asked to submit a solution s such that
 - E.g., $H(r||s) \leq t = 2^{L-c}$
 - If solution is invalid, then server ignores request r

Defer DDoS with PoW

- Use PoW to defer Denial-of-Service (DDoS)
 - Server uses counter c as difficulty d of PoW
- No. of hashes (i.e., time) to find a solution for the next request is 2^c , increases **exponentially** with no. of requests
 - Attacker cannot submit a **large** number of requests within a short time.
 - Only **minor** impacts to legitimate users

- Example: UC server goes down if receives >30 requests within 1 min from one user/attacker
- User/attacker needs to compute 2^c hashes for each new request, c is counter
- **Minor impact** on Alice (200,000 hashes/sec)
 - 1st request, no hash, no extra time;
 - 2nd request, 2 hashes, 1 micro sec;
 - 3rd request, 2^2 hashes, 2 micro secs;
 - 4th request, 2^3 hashes, 4 micro secs;
 - 5th request, 2^4 hashes, 8 micro secs;
- Total overhead: only 15 micro secs

- Example: UC server goes down if receives >30 requests within 1 min from one user/attacker
- User/attacker needs to compute 2^c hashes for each new request, c is counter
- **Defer DDoS attack** (200,000 hashes/sec)
 - First 20 requests, 2^{20} hashes in total, 5 secs;
 - 21st request, 2^{20} hashes, 5 secs;
 - 22nd request, 2^{21} hashes, 11 secs;
 - 23rd request, 2^{22} hashes, 21 secs;
 - 24th request, 2^{23} hashes, **42** secs;
- Attacker cannot submit >30 requests within 1min

DDoS Attack

- With PoW, DDoS attack may still work if
 - Attacker has **enormous** hashing power



- Example: UC server goes down if receives >30 requests within 1 min from one user/attacker
- User/attacker needs to compute 2^c hashes for each new request, c is counter
- DDoS attack ($200,000,000,000$ hashes/sec)
 - 21st request, 2^{20} hashes, 5 micro secs;
 - 22nd request, 2^{21} hashes, 11 micro secs;
 - 23rd request, 2^{22} hashes, 21 micro secs;
 - 24th request, 2^{23} hashes, 42 micro secs;
- Attacker can submit 30 requests within 1 sec, finish DDoS attack easily

DDoS Attack

- With PoW, DDoS attack may still work if
 - Attacker controls a **huge** number of machines
- Example:
 - Attacker controls **1,000** machines
 - Counter is still 1 for each machine
 - PoW is not effective yet
 - **1,000** requests arrive same time at UC server
 - UC server goes down again...

Bitcoin Mining


- Adding new blocks to blockchain is called [mining](#)
- There are many miners in the network, how to decide which miner generates the next block?
 - Each miner is given a same “Puzzle”
 - Each miner selects its own block from current pending trans pool
 - Who finds a solution for “Puzzle” first based on its block will add its block to blockchain

PoW in Bitcoin

- Each miner is given a same PoW target
- Who finds a PoW solution first based on its block will add its block to blockchain.
- This PoW solution can be verified by other miners, once verified, this block is added to the blockchain, everyone will start to find a solution for the next block.

- Example: Alice and Bob are two miners
- There are 10,000 pending trans in the network
 - Alice has 8,000 pending trans locally
 - Bob has 9,000 pending trans locally
 - Some trans are not reached Alice or Bob yet
- Block **No. 12345** is just added to blockchain, the current target is **$t = 00001111$** for all miners
 - All miners start to create next block
 - Alice forms a block with 2,000 trans (about 1MB)
 - Bob forms a block with 1,800 trans (about 1MB)
 - Alice's block and Bob's block are different

- Example: Alice and Bob are two miners
 - Current block length: 12345,
 - Blockchain hash value: $h=0xabcd$
 - Current PoW target: $t = 00001111$
-
- Alice forms a block with 2,000 trans (about 1MB)
 - Alice builds a Merkle tree based on those 2,000 trans, computes a **root hash** $r=0xdcef$



- Example: Alice and Bob are two miners
- Current block length: 12345,
- Blockchain hash value: $h=0xabcd$
- Current PoW target: $t = 00001111$
- Given $h=0xabcd$, $r=0xdcef$, and target t , Alice starts to find a **solution s** , s.t.
$$H(h \parallel r \parallel s) \leq t$$
- If Alice finds a solution before others, she broadcasts her block with (r, s) .

- Example: Alice and Bob are two miners
- Current block length: 12345,
- Blockchain hash value: $h=0xabcd$
- Current PoW target: $t = 00001111$

- Others verify (r, s) , with current (h, t) , if this (r, s) passes, this block with the 2,000 trans selected by Alice will be added as block [No. 12346](#).

- Update hash value as $h = H(h \parallel r \parallel s)$
- Block No. 12346 & new h are broadcasted
- All miners start to build next block

- Example: Alice and Bob are two miners
- Current block length: 12345,
- Blockchain hash value: $h=0xabcd$
- Current PoW target: $t = 00001111$

- Alice and Bob are both finding a PoW solution
- In PoW, enumerating all possible nonces is the only way to find a solution

- Who finds a solution first depends hashing power
 - Alice 200 hashes/sec; Bob 100 hashes/sec
 - In theory, Alice will find a solution first

- Example: Alice and Bob are two miners
- Current PoW target: $t = 00001111$; difficulty is 4
- Average $2^4 = 16$ hashes to find a solution
- Alice 200 hashes/sec; Bob 100 hashes/sec
- If Alice & Bob choose exactly the same block
 - same set of trans, same order of trans
- Alice & Bob have exactly same PoW problem
- If Alice & Bob use same approach to find solution
 - Same initial nonce, same enumeration (n++)
 - E.g., both need 14 hashes to find solution
- Then, Alice will definitely win

- Example: Alice and Bob are two miners
- Current PoW target: $t = 00001111$; difficulty is 4
- Average $2^4 = 16$ hashes to find a solution
- Alice 200 hashes/sec; Bob 100 hashes/sec
- Alice & Bob have different pending trans
- Alice & Bob select 2 different blocks
 - different sets of trans, different orders of trans
- Same target, but 2 different PoW problems
- Different enumeration methods
 - E.g., Alice: $n=n+1$; Bob: $n=n+2$
- Both could win; Alice has a better chance

- Example: Alice and Bob are two miners
- Current PoW target: $t = 00001111$; difficulty is 4
- Average $2^4 = 16$ hashes to find a solution
- Alice 200 hashes/sec; Bob 100 hashes/sec
- On average, both need 16 hashes to find solution
- Bob could still win sometimes

Block	Alice's Hashes	Alice's Time	Bob's Hashes	Bob's Time
12345	22	0.11 s	10	0.10 s
12346	15	0.075 s	20	0.20 s
12347	11	0.055 s	17	0.17 s

- Example: Alice and Bob are two miners
- Current PoW target: $t = 00001111$; difficulty is 4
- Average $2^4 = 16$ hashes to find a solution
- Alice 200 hashes/sec; Bob 100 hashes/sec
- Total hashing power in Bitcoin network
 - $T = 200 + 100 = 300$ hashes/sec
- Outputs of hash func. are uniformly distributed
- Probability that Alice can win
 - $p_A = 200/T = 2/3$
- Probability that Bob can win
 - $p_B = 100/T = 1/3$

- Example: Alice and Bob are two miners
- Alice 200 hashes/sec; Bob 100 hashes/sec
 - Total hashing power in Bitcoin network
 - $T = 200 + 100 = 300$ hashes/sec
 - Probability that Alice can win
 - $p_A = 200/T = 2/3 = 0.67$
- Alice increases to 400 hashes/sec, Bob is the same
 - Total hashing power in Bitcoin network
 - $T = 400 + 100 = 500$
 - Probability that Alice can win
 - $p_A = 400/T = 4/5 = 0.80$

- Example: Alice and Bob are two miners
- Alice 200 hashes/sec; Bob 100 hashes/sec
 - Total hashing power in Bitcoin network
 - $T = 200 + 100 = 300$ hashes/sec
 - Probability that Alice can win
 - $p_A = 200/T = 2/3 = 0.67$
- Alice 400 hashes/sec, Bob 200 hashes/sec
 - Total hashing power in Bitcoin network
 - $T = 400 + 200 = 600$
 - Probability that Alice can win
 - $p_A = 400/T = 2/3 = 0.67$

Mining Competition

- Every miner realizes that increasing hashing power could increase its probability to win
- Hashing power competition
 - CPU (2011, 2012) —> GPU (2013) —> Hardware (2015) —> Mining Station (2016, 2017)



AntMiner S7

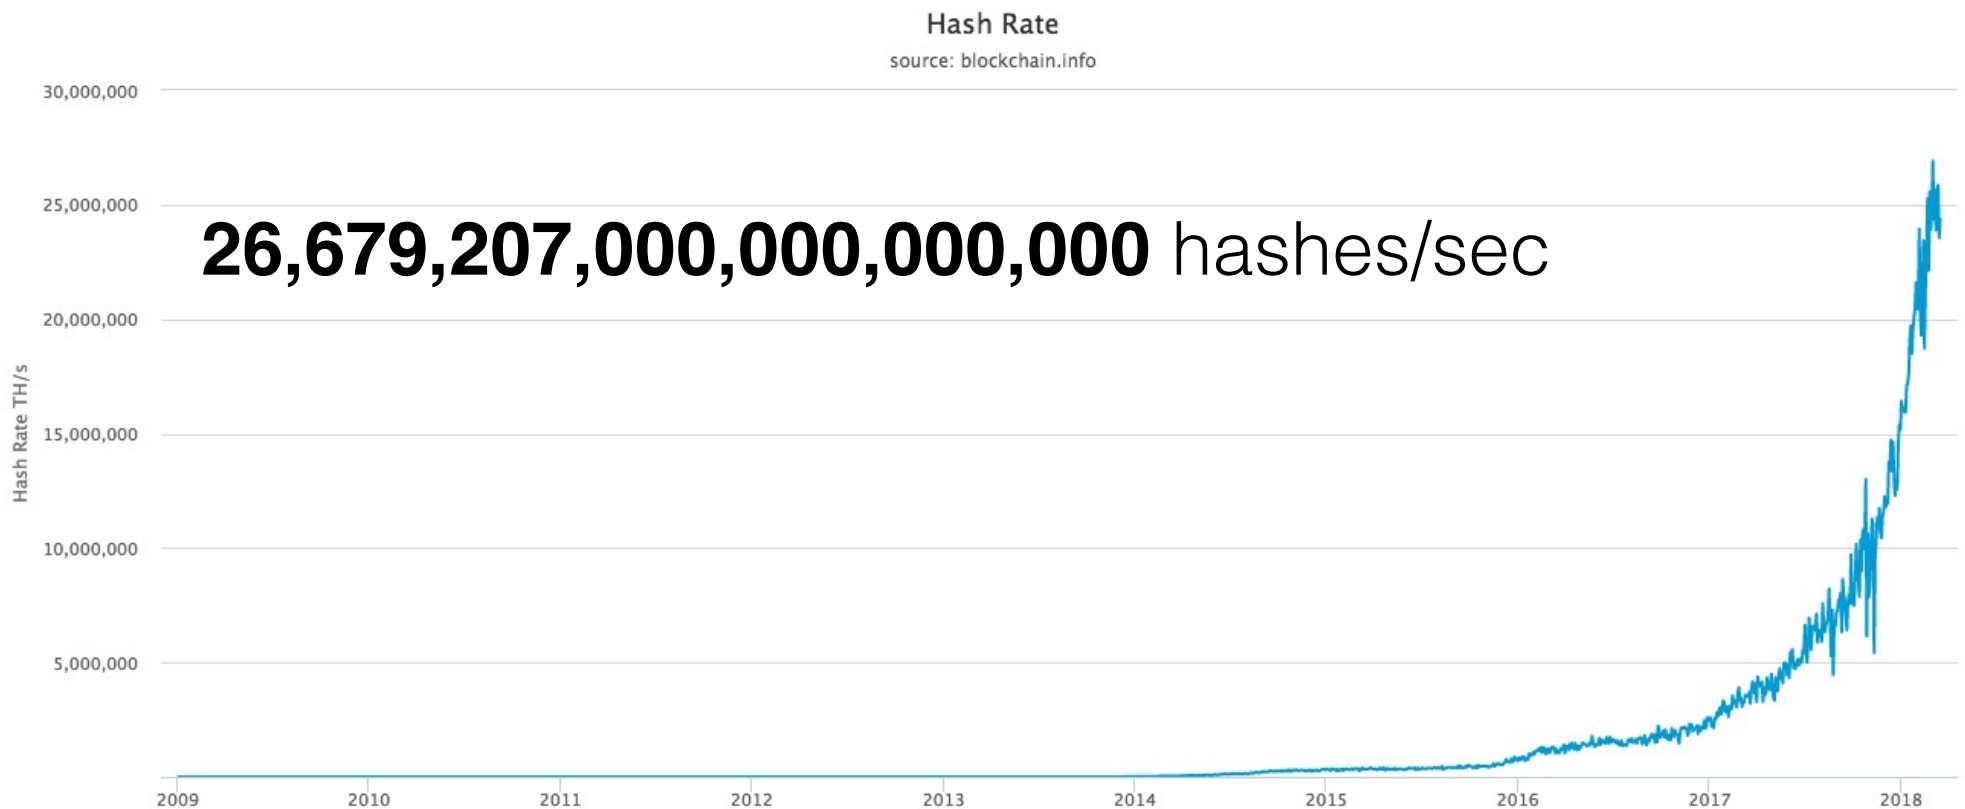


Avalon6



Mining Competition

- All users increase hashing power
- The total hashing power increases **exponentially**



Mining Competition

- If you buy a \$2,000 PC to mine bitcoins yourself
 - 400,000 hashes/sec
 - About 1×10^{-15}
- But other miners are like this...



Mining Competition

- You weapons are like this...



- Other miners' weapons are like this..



Mining Group

- New Hope: Individual miners form mining groups in order to have a better chance to win.
- Alice, Bob and Charlie form a mining group
- Alice: 400 hashes/sec; Bob: 200 hashes/sec;
Charlie: 1000 hashes/sec;
- Group: 1600 hashes/sec
- A mining group may have several thousand miners

Mining Group

- Members in a group try to solve a **same PoW** problem with different subsets of randoms.
- If one member finds a solution, this group gets transaction fee and new bitcoins
- All members **share** this reward based on hashing power: Alice: 0.25; Bob 0.125; Charlie: 0.625
- **David** who organizes this group will get some reward as well

Mining Group

